

ME EN 7960: Scientific Machine Learning Fall 2024

Final Project: Operator Learning on Curved Geometries:
DeepONet Architectures for Poisson Problems on the Torus

Milena Belianovich
April 28, 2025

1 Overview and Motivation

The goal of this project is to investigate how Deep Operator Networks (DeepONets) can be adapted to learn solution operators of partial differential equations (PDEs) defined on curved manifolds, specifically the torus. This extends the traditional application of DeepONets, which are typically designed for Euclidean domains, and tests their ability to generalize on non-Euclidean geometries.

Learning operators on manifolds has significant relevance for scientific computing and engineering applications, such as modeling physical phenomena on curved surfaces (e.g., geophysics, fluid flows over curved surfaces, and material science problems). Our focus is on the Poisson equation on a torus, for which we generate synthetic datasets using high-accuracy numerical solvers and study various architectural adaptations of DeepONets to improve performance.

1.1 Problem Setup

We study the Poisson equation

$$-\Delta_{\mathcal{M}} u(\mathbf{x}) = f(\mathbf{x})$$

on the 2D torus embedded in \mathbb{R}^3 , aiming to learn the operator $f \mapsto u$ from forcing terms to solutions using data-driven neural network models.

2 Data Generation

The first phase of the project was the generation of training and testing data for learning solution operators on manifolds. Specifically, we solved the Poisson equation on a torus using a radial basis function finite difference (RBF-FD) method.

Originally, the PDE solver was implemented in MATLAB through the provided `DriverPoissonSolveTorus.m` code. To fully integrate data generation into our PyTorch training pipeline, we rewrote the solver in Python. This was achieved through the `data_gen.py` script, which solves the Poisson equation with manufactured solutions, samples the toroidal

mesh points, computes corresponding forcing terms, and saves the datasets in `.npz` format for later training.

The PDE being solved was the Poisson equation on a torus surface:

$$-\Delta_{\mathcal{M}} u(\mathbf{x}) = f(\mathbf{x}), \quad \mathbf{x} \in \text{Torus}$$

where $\Delta_{\mathcal{M}}$ denotes the Laplace–Beltrami operator on the manifold.

Testing Different Manufactured Solutions (ξ Values). The true solution was chosen as:

$$u(\mathbf{x}) = \sin(\xi\theta) \sin(\xi\phi)$$

where θ, ϕ are the intrinsic toroidal coordinates and ξ is a frequency-like parameter controlling the oscillation complexity of the solution.

We experimented with several values of $\xi \in \{2, 4, 6, 8\}$ to analyze their effect on solver accuracy and learning complexity. Figures 1 show representative plots comparing the true and numerical solutions for different ξ and mesh sizes.

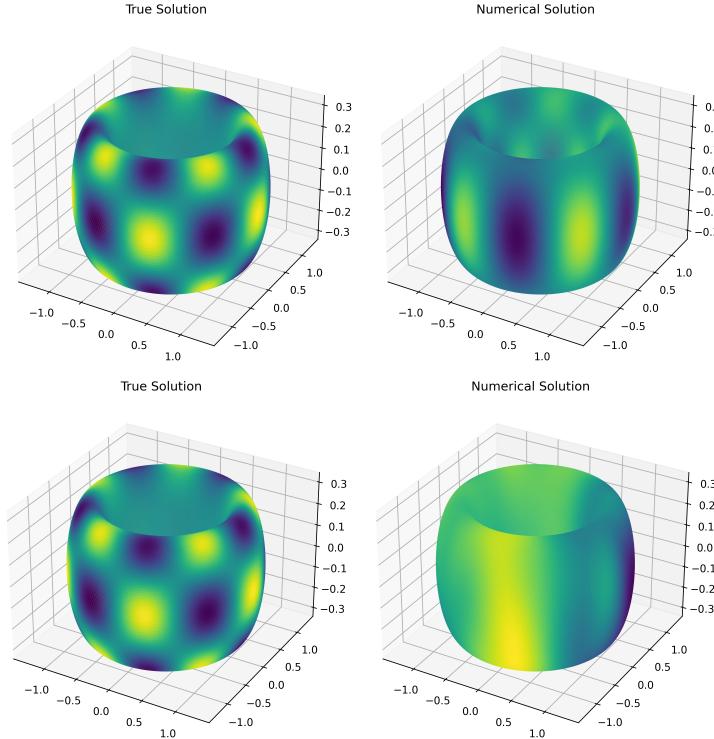


Figure 1: Comparison of true and numerical solutions on the torus for $\xi = 2$ (top) and $\xi = 4$ (bottom) at high mesh resolution.

We observed that for very large ξ (e.g., $\xi = 8$), the solver struggled to accurately resolve the solution, leading to high numerical errors even at fine discretizations. Lower ξ values (such as $\xi = 2$) produced overly smooth solutions that were less challenging for operator learning. Based on both visual inspection and solver error analysis, we ultimately selected $\xi = 4$ as the manufactured solution used in all subsequent DeepONet training.

Solver Convergence. The RBF-FD solver convergence was validated by plotting the relative L^2 errors against mesh resolution for different ξ values, as shown in Figure 2.

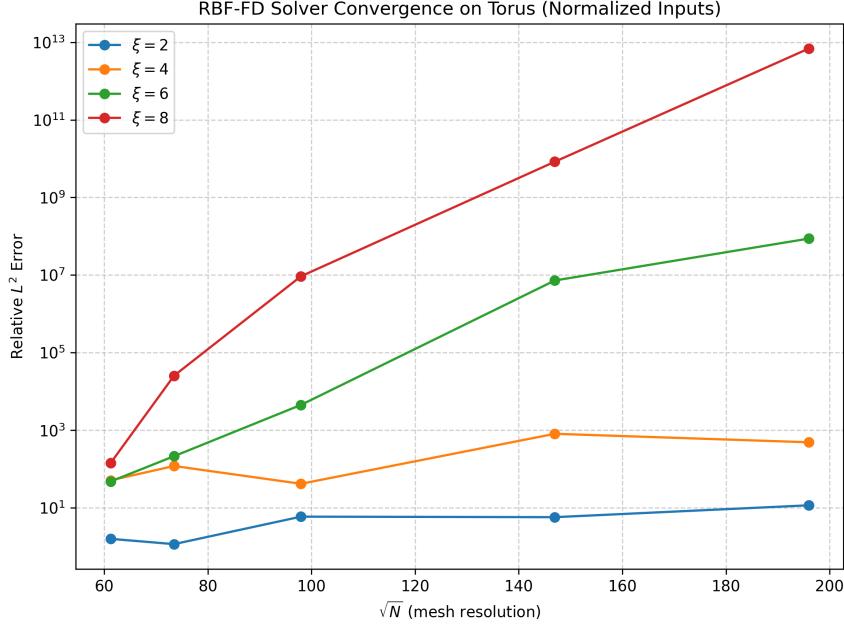


Figure 2: RBF-FD solver convergence plot for different ξ values on the torus.

The convergence plots highlight that:

- $\xi = 2$ achieved the lowest relative error, but produced solutions that were too smooth.
- $\xi = 4$ maintained reasonable solver error while offering enough complexity to meaningfully train neural operators.
- $\xi = 6$ and $\xi = 8$ resulted in diverging numerical errors, particularly as mesh resolution increased, indicating numerical instability or excessive stiffness.

Based on these observations, $\xi = 4$ was chosen for all datasets used in model training and evaluation.

3 DeepONet Architectures

We investigate several variations of DeepONet architectures to study their operator learning capabilities:

3.1 Vanilla DeepONet

The standard DeepONet architecture uses a **branch network** that encodes the forcing term f and a **trunk network** that encodes the spatial coordinates x .

The prediction is given by:

$$u(x) = \langle \text{Branch}(f), \text{Trunk}(x) \rangle,$$

where $\langle \cdot, \cdot \rangle$ denotes an inner product in the latent space.

The vanilla DeepONet serves as the baseline. However, because the problem domain is a curved manifold, we hypothesize that incorporating geometric information such as surface normals may improve predictive accuracy. We therefore introduce three extended architectures that leverage additional input modalities and nonlinear fusion strategies.

3.2 Dual-Branch DeepONet (with Normals)

To better incorporate the geometric structure of the torus, we extend the vanilla architecture by adding an additional branch that inputs the surface normals $n(x)$.

The new prediction formula becomes:

$$u(x) = \langle \text{Branch}_f(f), \text{Trunk}(x) \rangle + \langle \text{Branch}_n(n(x)), \text{Trunk}(x) \rangle.$$

3.3 Fusion DeepONet (Nonlinear Combination)

We further modify the architecture by fusing the two branch outputs via a nonlinear multi-layer perceptron (MLP) before producing the final prediction.

Given:

$$\text{FusionInput}(x) = [\langle \text{Branch}_f(f), \text{Trunk}(x) \rangle, \langle \text{Branch}_n(n(x)), \text{Trunk}(x) \rangle],$$

the final output is:

$$u(x) = \text{FusionMLP}(\text{FusionInput}(x)),$$

where the FusionMLP introduces additional nonlinearity.

3.4 Fusion2 DeepONet (Branch Combination)

In an alternative design, we concatenate the inputs f and $n(x)$ at the branch level and pass them through a single branch network:

$$u(x) = \langle \text{Branch}(f, n(x)), \text{Trunk}(x) \rangle.$$

4 Training Setup

Training is performed using the Adam optimizer with cosine learning rate decay. A batch size of 512 is used for all models, and the learning rate is scheduled to decay smoothly during training.

Each architecture is trained to minimize the mean squared error (MSE) loss between predicted and true solution values.

We split the dataset into:

- Training set (majority of points),
- Test set (a small percentage held out for monitoring performance),

- Generalization set (an additional random subset not seen during training, to compute final generalization error).

Training is performed using the Adam optimizer with cosine learning rate decay. All models use hidden layers of size [128, 128] unless otherwise specified.

5 Evaluation Metrics

We track:

- Training error (relative L^2 norm, normalized by true solution norm),
- Test error (relative L^2 norm, normalized),
- Generalization error (computed on unseen test points, normalized).

Plots of training and testing loss versus epoch are generated for each architecture and dataset size.

6 Results and Analysis

6.1 Vanilla DeepONet Baseline

We first trained the vanilla DeepONet architecture to serve as a baseline. Training, testing, and generalization errors were tracked across different mesh resolutions $N \in \{3750, 5400, 9600, 21600, 38400\}$.

Figure 3 shows the training and testing loss curves for $N = 9600$, along with the final generalization error.

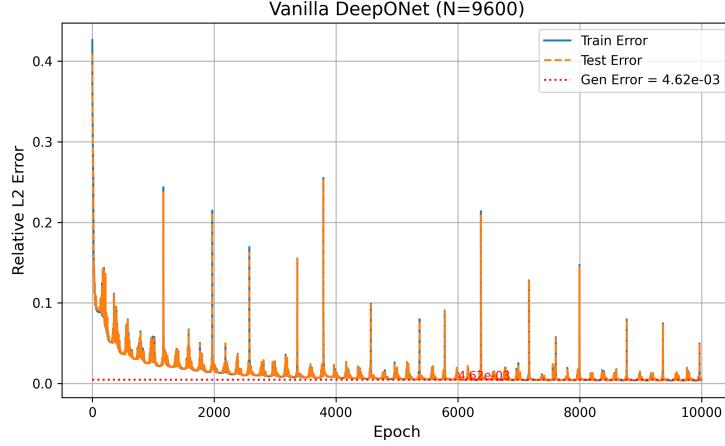


Figure 3: Training and testing errors for Vanilla DeepONet ($N = 9600$).

The predicted solution compared to the true solution for $N = 9600$ is shown in Figure 4.

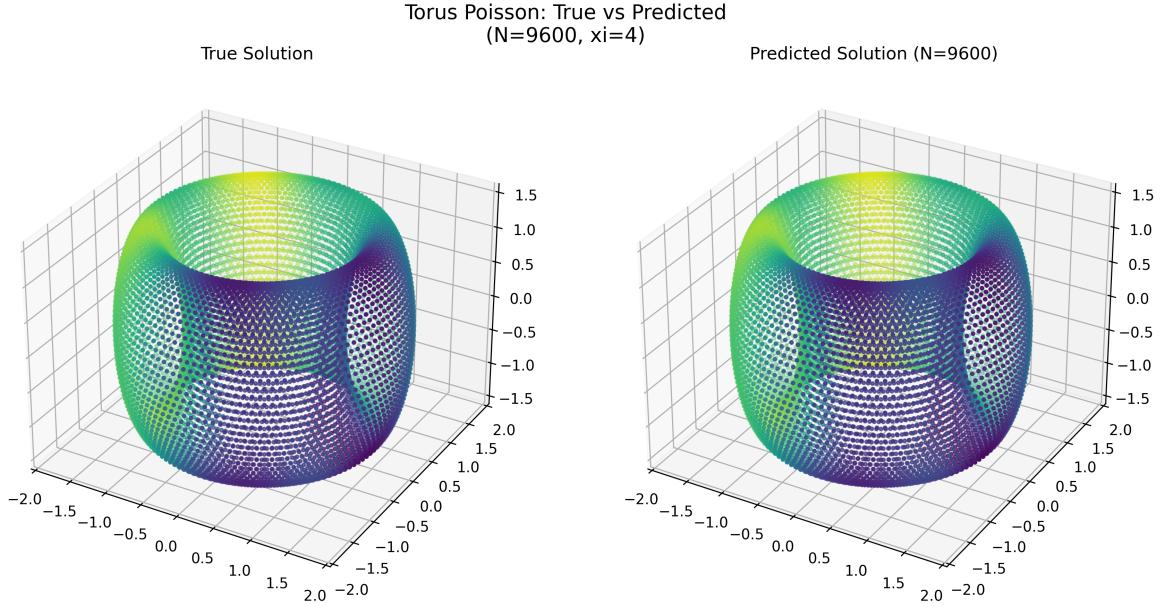


Figure 4: True vs predicted solutions for Vanilla DeepONet ($N = 9600$).

The vanilla DeepONet was able to learn coarse features of the solution but exhibited relatively high generalization errors compared to the extended architectures introduced later.

6.2 New DeepONet (Dual Branch)

The New DeepONet architecture added an additional branch network to encode surface normals $n(x)$ alongside the forcing term f .

Figure 5 shows the training and testing loss curves for $N = 9600$.

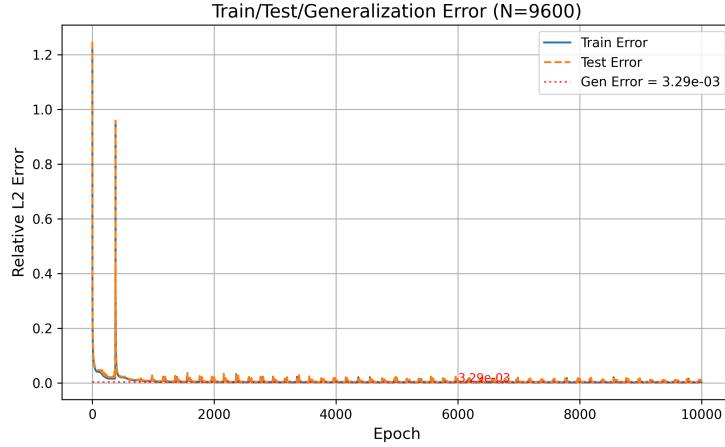


Figure 5: Training and testing errors for New DeepONet (Dual Branch) ($N = 9600$).

Figure 6 shows the true versus predicted solution comparison.

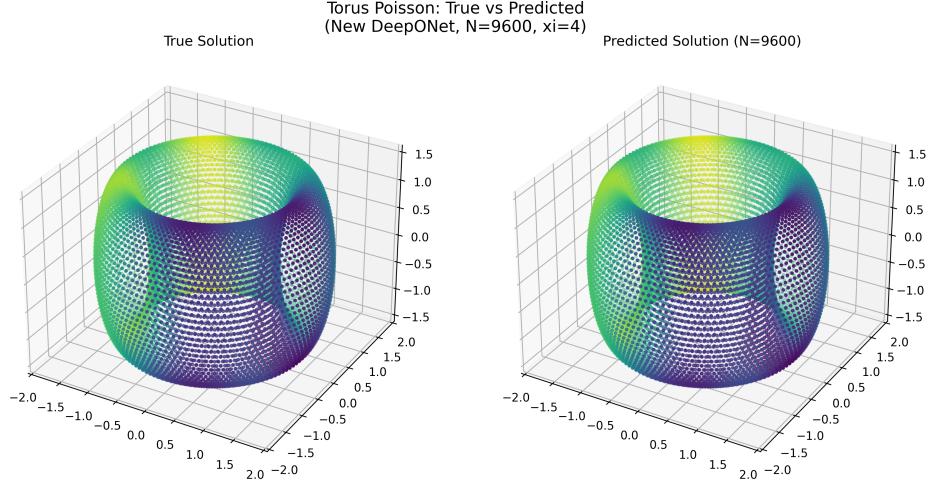


Figure 6: True vs predicted solutions for New DeepONet (Dual Branch) ($N = 9600$).

We observe that the dual-branch structure improved test and generalization errors compared to vanilla, although gains were modest in some cases.

6.3 Fusion DeepONet (Nonlinear Combination)

In the Fusion DeepONet, the outputs of the two branches were combined using a nonlinear MLP (FusionMLP).

Training and testing errors for $N = 9600$ are shown in Figure 7.

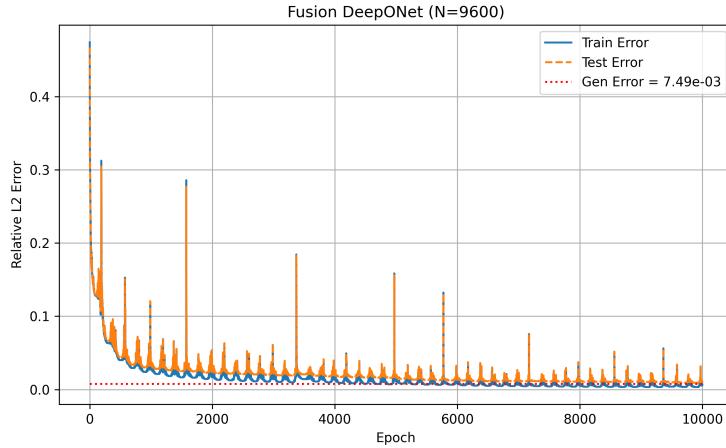


Figure 7: Training and testing errors for Fusion DeepONet ($N = 9600$).

Figure 8 shows the prediction comparison.

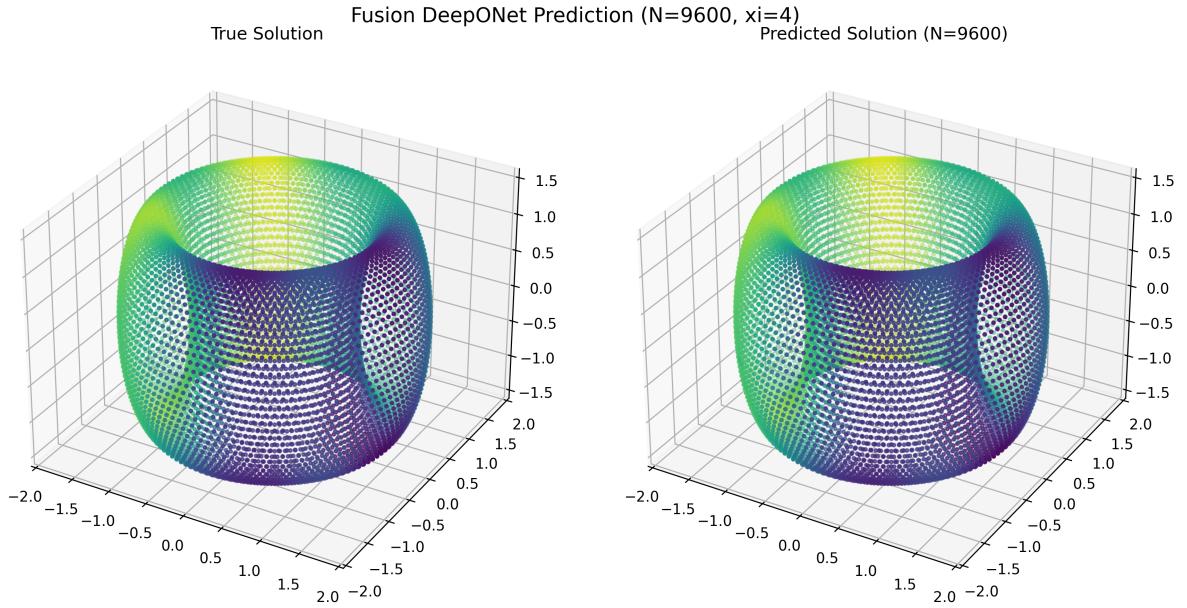


Figure 8: True vs predicted solutions for Fusion DeepONet ($N = 9600$).

Fusion DeepONet provided improvement over the vanilla baseline and occasionally achieved competitive performance relative to Fusion2, especially at intermediate mesh sizes.

6.4 Fusion2 DeepONet (Branch Combination)

In Fusion2 DeepONet, the forcing term f and normals $n(x)$ were concatenated at the input level before passing through a single branch network.

Training and testing losses for $N = 9600$ are shown in Figure 9.

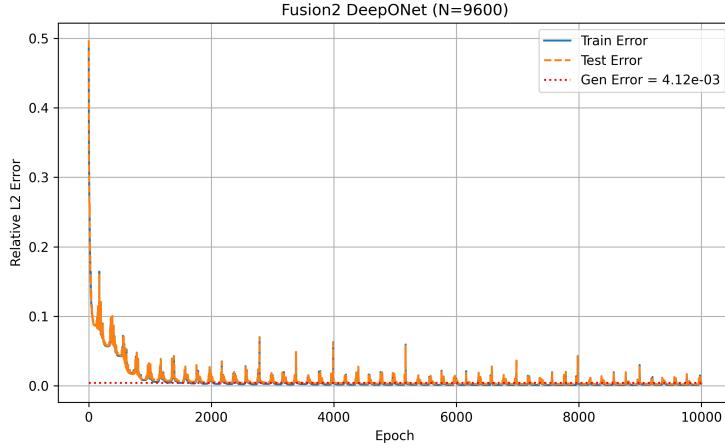


Figure 9: Training and testing errors for Fusion2 DeepONet ($N = 9600$).

Figure 10 shows the true vs predicted solutions.

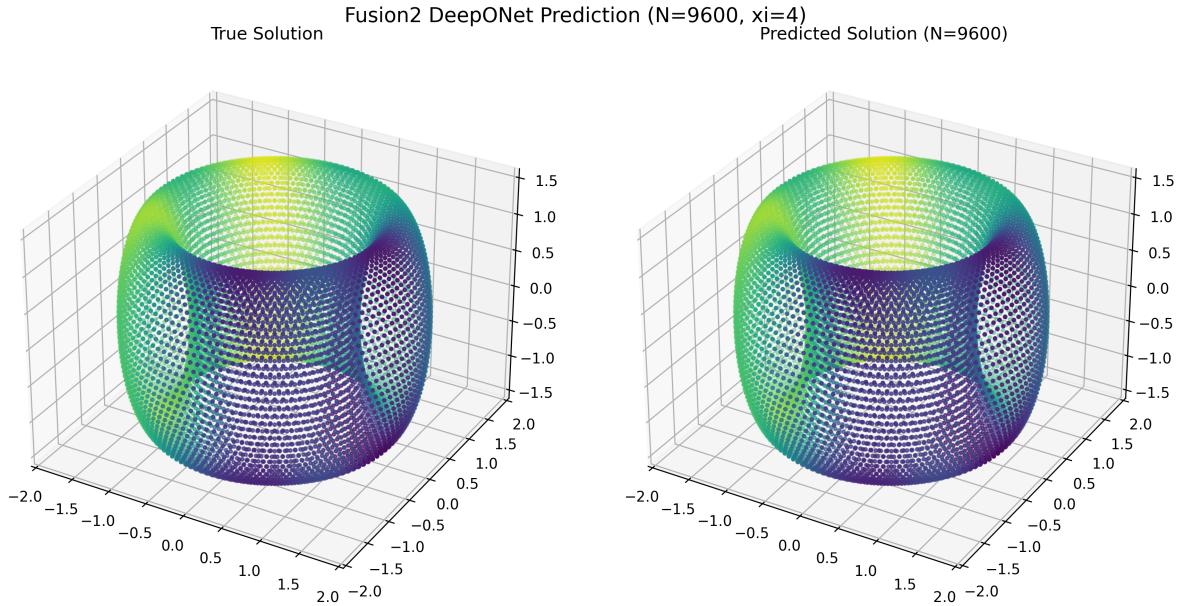


Figure 10: True vs predicted solutions for Fusion2 DeepONet ($N = 9600$).

Fusion2 consistently achieved the lowest generalization errors across all N values tested, indicating that early-stage combination of the forcing and geometric information is beneficial.

6.5 Comparison Across Architectures

To directly compare the behavior of all architectures, we generated combined comparison plots. Figure 11 shows training and testing losses across all architectures for $N = 9600$.

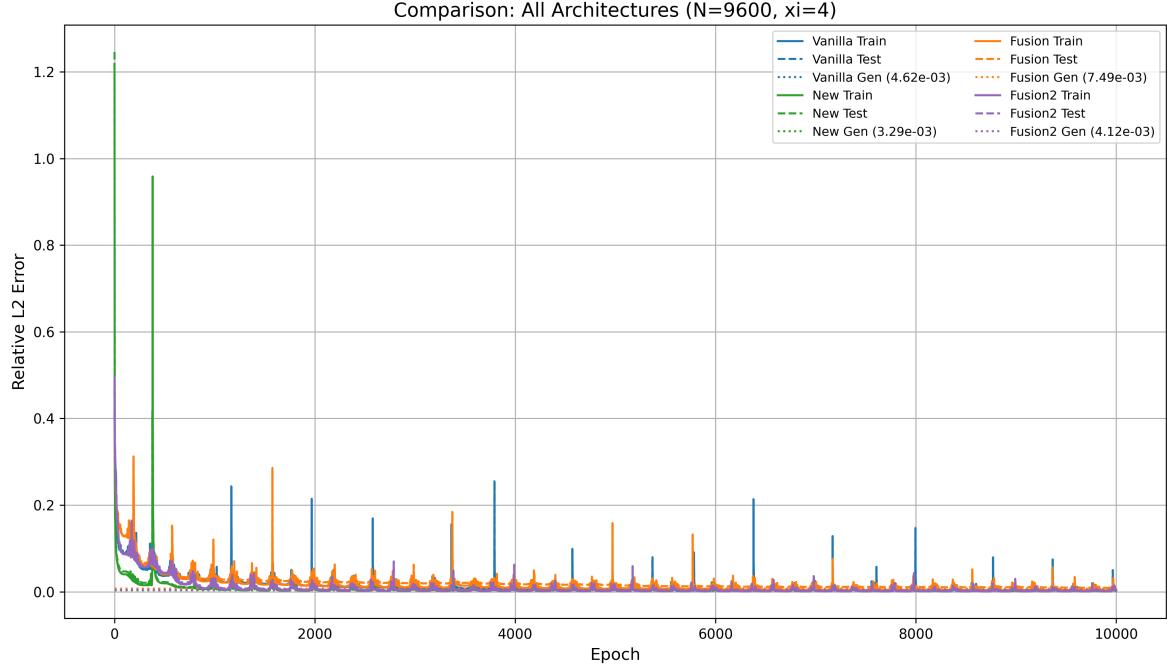


Figure 11: Comparison of training, testing, and generalization errors for all DeepONet architectures ($N = 9600$).

We note:

- Vanilla DeepONet had the highest training and generalization errors.
- The dual-branch New DeepONet achieved lower errors, validating the importance of geometric information.
- Fusion DeepONet improved on vanilla for most N , and for some cases ($N = 3750$, $N = 5400$) matched or slightly outperformed New DeepONet.
- Fusion2 DeepONet consistently achieved the lowest or near-lowest generalization errors across all mesh sizes.

A quantitative summary is provided in Table 1:

N	Vanilla Gen Error	New Gen Error	Fusion Gen Error	Fusion2 Gen Error
3750	7.2390e-03	4.5530e-03	3.8511e-03	4.6584e-03
5400	5.2175e-03	3.4584e-03	9.3149e-03	3.5701e-03
9600	4.6211e-03	3.2863e-03	7.4857e-03	4.1249e-03
21600	5.0316e-03	2.5856e-03	3.9297e-03	3.6781e-03
38400	1.7675e-02	1.8749e-02	1.8701e-02	1.7457e-02

Table 1: Generalization errors across architectures.

6.6 Discussion

This study demonstrates that careful integration of geometric information and nonlinear fusion strategies significantly improves the operator learning performance of DeepONets on curved domains. Fusion2 DeepONet, in particular, offers a promising architecture for learning PDE solution operators on manifolds.

Further improvements could be achieved by deeper networks, spectral normalization, or using different function bases for branch/trunk embeddings.